

```
In [1]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import numpy as np
import pandas as pd
import tensorflow as tf
from tensorflow import keras
from sklearn.model_selection import train_test_split, StratifiedKFold
from sklearn.preprocessing import StandardScaler
from sklearn.utils.class_weight import compute_class_weight
from sklearn.metrics import accuracy_score, precision_score, recall_score,
import matplotlib.pyplot as plt
import seaborn as sns
```

```
/Users/mengjiayu/opt/anaconda3/lib/python3.8/site-packages/pandas/core/computation/expressions.py:20: UserWarning: Pandas requires version '2.7.3' or newer of 'numexpr' (version '2.7.1' currently installed).
  from pandas.core.computation.check import NUMEXPR_INSTALLED
```

```
In [2]: df = pd.read_csv('output_dataset.csv')
```

```
In [4]: whole = df[['Total Bilirubin', 'Alkphos Alkaline Phosphotase',
'Sgpt Alamine Aminotransferase', 'Sgot Aspartate Aminotransferase',
'Total Protiens', 'ALB Albumin', 'A/G Ratio Albumin and Globulin Ratio',
'Bilirubin Ratio', 'SGOT/SGPT Ratio', 'Protien Ratio', 'Result']]
```

```
In [5]: whole
```

```
Out[5]:
```

	Total Bilirubin	Alkphos Alkaline Phosphotase	Sgpt Alamine Aminotransferase	Sgot Aspartate Aminotransferase	Total Protiens	ALB Albumin	Alk Glc
0	0.530628	187.0	2.833213	2.944439	6.8	3.3	
1	2.476538	690.0	4.174387	4.615121	7.5	3.2	
2	2.116256	490.0	4.110874	4.234107	7.0	3.3	
3	0.693147	182.0	2.772589	3.044522	6.8	3.4	
4	1.589235	195.0	3.332205	4.094345	7.3	2.4	
...	
30686	1.163151	610.0	2.890372	3.367296	7.3	2.6	
30687	1.360977	482.0	3.135494	3.555348	7.0	2.4	
30688	2.054124	542.0	4.762174	4.204693	6.4	3.1	
30689	1.064711	231.0	2.833213	4.025352	4.6	1.8	
30690	1.410987	253.0	4.394449	6.007239	6.8	3.9	

30691 rows × 11 columns

```
In [ ]:
```

```
In [ ]:
```

```

In [6]: # Prepare features and target variable
feature_cols = ['Total Bilirubin', 'Alkphos Alkaline Phosphotase',
                'Sgpt Alamine Aminotransferase', 'Sgot Aspartate Aminotrans',
                'Total Protiens', 'ALB Albumin', 'A/G Ratio Albumin and Glob',
                'Bilirubin Ratio', 'SGOT/SGPT Ratio', 'Protien Ratio']

X = df[feature_cols]
y = df['Result']

# Convert target labels: 2 → 0 (Non-Disease), 1 → 1 (Disease)
y = (y == 1).astype(int)

# Train-test split (80:20)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran

# Standardize Numeric Features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Compute Correct Class Weights
class_weights = compute_class_weight('balanced', classes=np.array([0, 1]), y
class_weights_dict = {0: class_weights[0], 1: class_weights[1]} # Ensure or

# Function to Build FNN Model
def build_fnn_model(hidden_units=128, dropout_rate=0.3, learning_rate=0.001):
    model = keras.Sequential([
        keras.layers.Dense(hidden_units, activation='relu', input_shape=(X_train.shape[1],)),
        keras.layers.BatchNormalization(),
        keras.layers.Dropout(dropout_rate),
        keras.layers.Dense(hidden_units // 2, activation='relu'),
        keras.layers.Dense(hidden_units // 4, activation='relu'),
        keras.layers.Dense(1, activation='sigmoid') # Binary classification
    ])

    model.compile(optimizer=keras.optimizers.Adam(learning_rate=learning_rate),
                  loss='binary_crossentropy',
                  metrics=['accuracy'])

    return model

# Hyperparameter tuning with Stratified K-Fold Cross Validation
best_model = None
best_score = 0
kf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

for units in [64, 128, 256]: # Wider range of neurons
    for dropout in [0.1, 0.3, 0.5]: # More dropout variations
        for lr in [0.001, 0.0005, 0.0001]: # Test higher and lower learning rates
            fold_scores = []

            for train_idx, val_idx in kf.split(X_train, y_train):
                X_train_fold, X_val_fold = X_train[train_idx], X_train[val_idx]
                y_train_fold, y_val_fold = y_train.iloc[train_idx], y_train.iloc[val_idx]

                model = build_fnn_model(hidden_units=units, dropout_rate=dropout, learning_rate=lr)

                # Add Early Stopping to prevent overfitting
                early_stopping = keras.callbacks.EarlyStopping(monitor='val_loss', patience=10)

                history = model.fit(X_train_fold, y_train_fold, epochs=50, validation_data=(X_val_fold, y_val_fold),
                                  class_weight=class_weights_dict)

```

```
val_loss, val_acc = model.evaluate(X_val_fold, y_val_fold, val_loader)
fold_scores.append(val_acc)

avg_score = np.mean(fold_scores)
print(f"Units: {units}, Dropout: {dropout}, LR: {lr} -> Avg Val Acc: {avg_score}")

if avg_score > best_score:
    best_score = avg_score
    best_model = model # Save the best model dynamically

# Train the best model on full training data
print("\n Best Model Found - Training on Full Training Data...")
final_model = best_model

history = final_model.fit(X_train, y_train, epochs=50, batch_size=32, verbose=0,
                          callbacks=[keras.callbacks.EarlyStopping(monitor='val_loss', min_delta=0.001, patience=5)],
                          class_weight=class_weights_dict)
```

```

Units: 64, Dropout: 0.1, LR: 0.001 -> Avg Val Accuracy: 0.9695
Units: 64, Dropout: 0.1, LR: 0.0005 -> Avg Val Accuracy: 0.9744
Units: 64, Dropout: 0.1, LR: 0.0001 -> Avg Val Accuracy: 0.8991
Units: 64, Dropout: 0.3, LR: 0.001 -> Avg Val Accuracy: 0.9457
Units: 64, Dropout: 0.3, LR: 0.0005 -> Avg Val Accuracy: 0.9375
Units: 64, Dropout: 0.3, LR: 0.0001 -> Avg Val Accuracy: 0.8414
Units: 64, Dropout: 0.5, LR: 0.001 -> Avg Val Accuracy: 0.8733
Units: 64, Dropout: 0.5, LR: 0.0005 -> Avg Val Accuracy: 0.8753
Units: 64, Dropout: 0.5, LR: 0.0001 -> Avg Val Accuracy: 0.7628
Units: 128, Dropout: 0.1, LR: 0.001 -> Avg Val Accuracy: 0.9869
Units: 128, Dropout: 0.1, LR: 0.0005 -> Avg Val Accuracy: 0.9841
Units: 128, Dropout: 0.1, LR: 0.0001 -> Avg Val Accuracy: 0.9621
Units: 128, Dropout: 0.3, LR: 0.001 -> Avg Val Accuracy: 0.9845
Units: 128, Dropout: 0.3, LR: 0.0005 -> Avg Val Accuracy: 0.9810
Units: 128, Dropout: 0.3, LR: 0.0001 -> Avg Val Accuracy: 0.9359
Units: 128, Dropout: 0.5, LR: 0.001 -> Avg Val Accuracy: 0.9709
Units: 128, Dropout: 0.5, LR: 0.0005 -> Avg Val Accuracy: 0.9652
Units: 128, Dropout: 0.5, LR: 0.0001 -> Avg Val Accuracy: 0.8713
Units: 256, Dropout: 0.1, LR: 0.001 -> Avg Val Accuracy: 0.9846
Units: 256, Dropout: 0.1, LR: 0.0005 -> Avg Val Accuracy: 0.9873
Units: 256, Dropout: 0.1, LR: 0.0001 -> Avg Val Accuracy: 0.9862
Units: 256, Dropout: 0.3, LR: 0.001 -> Avg Val Accuracy: 0.9857
Units: 256, Dropout: 0.3, LR: 0.0005 -> Avg Val Accuracy: 0.9865
Units: 256, Dropout: 0.3, LR: 0.0001 -> Avg Val Accuracy: 0.9803
Units: 256, Dropout: 0.5, LR: 0.001 -> Avg Val Accuracy: 0.9845
Units: 256, Dropout: 0.5, LR: 0.0005 -> Avg Val Accuracy: 0.9815
Units: 256, Dropout: 0.5, LR: 0.0001 -> Avg Val Accuracy: 0.9583

```

Best Model Found - Training on Full Training Data...

Epoch 1/50

768/768 [=====] - 1s 2ms/step - loss: 0.0613 - accuracy: 0.9805 - val_loss: 0.0478 - val_accuracy: 0.9879

Epoch 2/50

768/768 [=====] - 1s 2ms/step - loss: 0.0527 - accuracy: 0.9828 - val_loss: 0.0529 - val_accuracy: 0.9860

Epoch 3/50

768/768 [=====] - 1s 2ms/step - loss: 0.0588 - accuracy: 0.9808 - val_loss: 0.0527 - val_accuracy: 0.9871

Epoch 4/50

768/768 [=====] - 1s 2ms/step - loss: 0.0503 - accuracy: 0.9835 - val_loss: 0.0550 - val_accuracy: 0.9847

Epoch 5/50

768/768 [=====] - 1s 2ms/step - loss: 0.0485 - accuracy: 0.9837 - val_loss: 0.0540 - val_accuracy: 0.9879

Epoch 6/50

768/768 [=====] - 1s 2ms/step - loss: 0.0421 - accuracy: 0.9860 - val_loss: 0.0631 - val_accuracy: 0.9871

```
In [ ]: # Units: 256, Dropout: 0.1, LR: 0.0005 -> Avg Val Accuracy: 0.9873
```

```
In [7]: final_model.summary()
```

Model: "sequential_99"

Layer (type)	Output Shape	Param #
dense_396 (Dense)	(None, 256)	2816
batch_normalization_99 (Batch Normalization)	(None, 256)	1024
dropout_99 (Dropout)	(None, 256)	0
dense_397 (Dense)	(None, 128)	32896
dense_398 (Dense)	(None, 64)	8256
dense_399 (Dense)	(None, 1)	65

=====
 Total params: 45057 (176.00 KB)
 Trainable params: 44545 (174.00 KB)
 Non-trainable params: 512 (2.00 KB)
 =====

```
In [8]: # Save the Final Model
final_model.save("fnn_model.h5")
```

/Users/mengjiayu/opt/anaconda3/lib/python3.8/site-packages/keras/src/engine/training.py:3000: UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')`.
 saving_api.save_model(

```
In [14]: # Get Final Training Accuracy (Last Epoch)
train_accuracy = history.history['accuracy'][-1]

# Get Mean Cross-Validation Accuracy from Best Model
cv_accuracy = best_score

accuracy_df = pd.DataFrame({
    "Metric": ["Train Accuracy", "Mean CV Accuracy"],
    "Value": [train_accuracy, cv_accuracy]
})

accuracy_df
```

```
Out[14]:
```

	Metric	Value
0	Train Accuracy	0.986030
1	Mean CV Accuracy	0.987251

```
In [9]: # Model Testing (Predictions)
y_pred_prob = final_model.predict(X_test) # Get probabilities
y_pred = (y_pred_prob > 0.5).astype(int) # Convert probabilities to class

# Model Evaluation
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred_prob)

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
```

```

# Print Evaluation Results in Tabular Form
metrics_df = pd.DataFrame({
    "Metric": ["Accuracy", "Precision", "Recall", "F1 Score", "ROC-AUC"],
    "Value": [accuracy, precision, recall, f1, roc_auc]
})
print("\n Model Evaluation Metrics:")
print(metrics_df)

# Classification Report
print("\n Classification Report:")
print(classification_report(y_test, y_pred))

# Plot Metrics Curves (Loss & Accuracy)
plt.figure(figsize=(12,5))
plt.subplot(1,2,1)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training vs. Validation Loss')
plt.legend()

plt.subplot(1,2,2)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Training vs. Validation Accuracy')
plt.legend()
plt.show()

# Confusion Matrix Heatmap
plt.figure(figsize=(6,5))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=['No', 'Yes'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

# Precision-Recall Curve
precision_vals, recall_vals, _ = precision_recall_curve(y_test, y_pred_prob)

plt.figure(figsize=(6,5))
plt.plot(recall_vals, precision_vals, marker='.', label="FNN Model")
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.title("Precision-Recall Curve")
plt.legend()
plt.grid()
plt.show()

# ROC Curve
fpr, tpr, _ = roc_curve(y_test, y_pred_prob)

plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, marker='.', label=f"FNN Model (AUC = {roc_auc:.3f})")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.legend()
plt.grid()
plt.show()

```




192/192 [=====] - 0s 713us/step

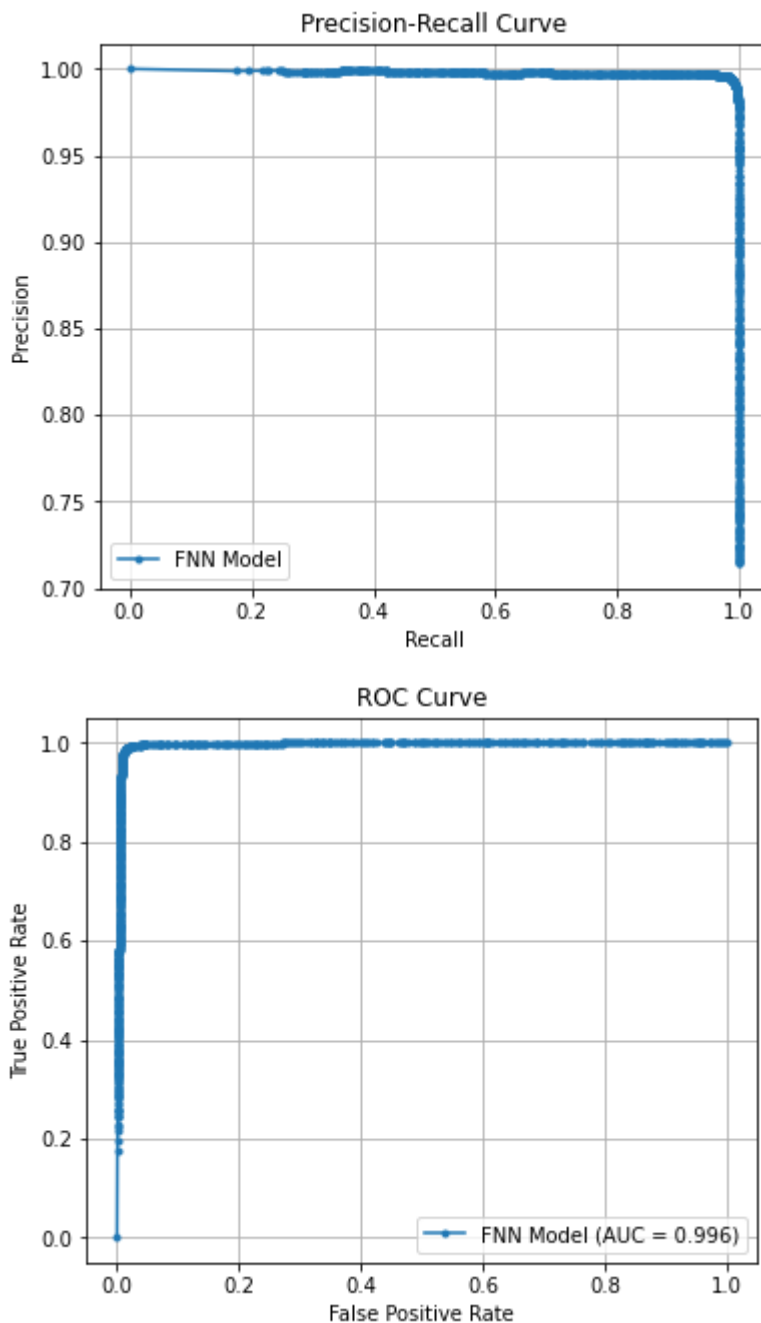
Model Evaluation Metrics:

	Metric	Value
0	Accuracy	0.987946
1	Precision	0.992234
2	Recall	0.990876
3	F1 Score	0.991554
4	ROC-AUC	0.996445

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.98	0.98	1755
1	0.99	0.99	0.99	4384
accuracy			0.99	6139
macro avg	0.98	0.99	0.99	6139
weighted avg	0.99	0.99	0.99	6139





```
In [15]: # Extract final training accuracy
train_accuracy = history.history['accuracy'][-1]
print(f"Final Training Accuracy: {train_accuracy:.4f}")
# Evaluate model on the test set
test_loss, test_accuracy = final_model.evaluate(X_test, y_test, verbose=0)
print(f"Test Accuracy: {test_accuracy:.4f}")
```

Final Training Accuracy: 0.9860
Test Accuracy: 0.9879

In []:

In []:

In []:

```
In [ ]: '''# Load the saved model
from tensorflow.keras.models import load_model

loaded_model = load_model("fnn_model.h5")
```



```
loaded_model.summary()  
'''
```